

# Tutorial 2.3 - Upper Confidence Bound

Tayyip Altan

October 2024

In this tutorial, we will be covering the Upper Confidence Bound (UCB) method to solve multi-armed bandit problems. We will be using the contextual package again. If you had trouble installing it, please visit the discussion board on Canvas to see additional hints.

## Dataset

For this tutorial, we will work with the Yahoo dataset, which was also used in tutorial 2.1.

```
# reads in full csv of Yahoo dataset
dfYahoo <- read.csv('yahoo_day1_10arms_tiny.csv')[,-c(1,2)]

# selects the two relevant columns from Yahoo dataset; arm shown to user and reward observed
dfYahoo_for_sim <- dfYahoo %>% select(arm, reward)
dfYahoo_for_sim$index <- 1:nrow(dfYahoo_for_sim)
```

## UCB: notation

As covered in the lecture, UCB methods decide on the arm to pick using an ‘optimistic’ estimate of the reward - e.g. which arm has the most potential to yield a high reward? This is done by picking the arm that has the highest combination of (1) the estimated reward,  $Q_t(a)$ , and (2) a bonus that rewards the arm for being underexplored, denoted as  $U_t(a)$ . The UCB algorithm picks the arm that maximizes the sum of these two factors:

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a) + c \cdot U_t(a). \quad (1)$$

(2)

In the general case of the UCB algorithm,  $U_t(a) = \sqrt{\frac{\log(t)}{N_t(a)}}$ , where  $N_t(a)$  is the number of times arm  $a$  has been pulled so far.

**Task 1: select the arm according to the UCB algorithm based on the data below.** Here, the code already provides a dataframe with per arm (1) the average reward, (2) the number of pulls thus far, and (3) an given value for  $t$ . Use this dataframe to determine which arm should be pulled if  $c = 0.1$ . The selected arm should be 8.

```

# set the seed
set.seed(0)

# get per item, the average reward and the number of items observed
dfYahoo_summary <- dfYahoo %>%
  group_by(arm) %>%
  summarise(avg_reward = mean(reward),
            n_pulls= n())

# the t in this case is simply the total of observations
t <- sum(dfYahoo_summary$n_pulls)
c <- 0.1

# TODO: Select an arm based on the UCB criterion (using equation (1))

```

## UCB: code

We will be simulating the performance of the UCB algorithm on the Yahoo! data, again with the contextual package.

```

# set the seed
set.seed(0)

## OfflineReplayEvaluatorBandit: simulates a bandit based on provided data
#
# Arguments:
#
#   data: dataframe that contains variables with (1) the reward and (2) the arms
#
#   formula: should consist of variable names in the dataframe. General structure is:
#             reward variable name ~ arm variable name
#
#   randomize: whether or not the bandit should receive the data in random order,
#              or as ordered in the dataframe.
#
# in our case, create a bandit for the data with 10 arms,
#   formula = reward ~ arm,
#   no randomization
bandit_yahoo <- OfflineReplayEvaluatorBandit$new(formula = reward ~ arm,
                                                data = dfYahoo,
                                                randomize = FALSE
                                                )

# lets generate a 10 simulations, each of size 10.000,
size_sim <- 10000
n_sim <- 10

# here we define the UCB policy object
# for some reason, the contextual package uses the 'UCB2Policy' tag for standard UCB algorithm
# in the package, 'alpha' is what we refer to as c

```

```

UCB_01      <- UCB2Policy$new(alpha=0.1)

# the contextual package works with 'agent' objects - which consist of a policy and a bandit
UCB_01_agent <- Agent$new(UCB_01, # add our UCB1 policy
                        bandit_yahoo) # add our bandit

## simulator: simulates a bandit + policy based on the provided data and parameters
#
# Arguments:
#
#   agent: the agent object, previously defined
#
#   horizon: how many observations from dataset used in the simulation
#
#   do_parallel: if True, runs simulation in parallel
#
#   simulations: how many simulations?
#

simulator    <- Simulator$new(UCB_01_agent, # set our agent
                             horizon= size_sim, # set the sizeof each simulation
                             do_parallel = TRUE, # run in parallel for speed
                             simulations = n_sim # simulate it n_sim times
                             )

# run the simulator object
history_UCB_01 <- simulator$run()

# gather results
df_Yahoo_UCB_01 <- history_UCB_01$data %>%
  select(t, sim, choice, reward, agent)

```

In the dataframe `df_Yahoo_UCB_01` we have gathered the results of the UCB policy for  $c = 0.1$ .

**Task 2:** repeat the steps of this tutorial, but now for  $c = 0.5$ .

```

# set the seed
set.seed(0)

# TODO: Adjust the code in this chunk to fit to the UCB policy for c = 0.5

# Define the policy
UCB_05      <- UCB2Policy$new(alpha=0.1)

# Define the agent
UCB_05_agent <- Agent$new(UCB_01, bandit_yahoo)

# Simulate
simulator    <- Simulator$new(list(UCB_01_agent, UCB_05_agent), # set our agents
                             horizon= size_sim, # set the sizeof each simulation
                             do_parallel = TRUE, # run in parallel for speed
                             simulations = n_sim, # simulate it n_sim times
                             )

```

```

# run the simulator object
history_UCB      <- simulator$run()

# gather results
df_Yahoo_UCB <- history_UCB$data %>%
  select(t, sim, choice, reward, agent)

```

**Task 3:** Make a plot that compares the UCB policy for  $c = 0.1$ ,  $c = 0.5$ . Compare the policies based on average cumulative reward. Can you conclude which one performs better, and if so why?

Your answer to Task 3 here:

```

# agent replaced with alpha value
df_Yahoo_UCB$agent <- recode(df_Yahoo_UCB$agent, 'UCB2' = '0.1', 'UCB2.2'='0.5')

# Maximum number of observations
max_obs <- 700

# data.frame aggregated for two agents
df_history_agg <- df_Yahoo_UCB %>%
  group_by(agent, sim)%>% # group by number of arms, the sim
  mutate(cumulative_reward = cumsum(reward))%>% # calculate cumulative sum
  group_by(agent, t) %>% # group by number of arms, the t
  summarise(avg_cumulative_reward = mean(cumulative_reward), # calc cumulative reward, se, CI
            se_cumulative_reward = sd(cumulative_reward, na.rm=TRUE)/sqrt(n_sim)) %>%
  mutate(cumulative_reward_lower_CI =avg_cumulative_reward - 1.96*se_cumulative_reward,
         cumulative_reward_upper_CI =avg_cumulative_reward + 1.96*se_cumulative_reward)%>%
  filter(t <=max_obs)

# TODO: make a plot using ggplot to compare UCB policy for c=0.1 and c=0.5
# 1: A plot that shows only the average cumulative rewards over time using the df_history_agg dataframe
# 2: The plot as defined in (1) together with the 95\% confidence interval.

```

**Task 4:** compare the UCB policy for  $c = 0.1$  to an  $\epsilon$ -greedy policy where  $\epsilon = 0.1$ , using the contextual package. Which one does better? Why do you think this might be?. Repeat this exercise again with *horizon* of 10.000.

```

# set the seed
set.seed(0)

# TODO: Add the EpsilonGreedyPolicy below to define the policy
eps_greedy <-

# TODO: Add the policy and the bandit below to define the agent
eps_greedy_agent <- Agent$new(, # Add the policy
                             ) # Add the bandit

# TODO: Finish the Simulator below
simulator <- Simulator$new(list(UCB_01_agent, eps_greedy_agent), # set our agents
                           horizon= , # set the sizeof each simulation
                           do_parallel = , # run in parallel for speed
                           simulations = , # simulate it n_sim times

```

```

)

# TODO: run the simulator object
history_compare_UCB_eps_greedy <-

# TODO: select relevant columns of the simulation
df_Yahoo_UCB_eps_greedy <-

# data.frame aggregated for two agents: UCB and epsilon greedy
df_history_agg <- df_Yahoo_UCB_eps_greedy %>%
  group_by(agent, sim)%>% # group by number of arms, the sim
  mutate(cumulative_reward = cumsum(reward))%>% # calculate cumulative sum
  group_by(agent, t) %>% # group by number of arms, the t
  summarise(avg_cumulative_reward = mean(cumulative_reward), # calc cumulative reward, se, CI
            se_cumulative_reward = sd(cumulative_reward, na.rm=TRUE)/sqrt(n_sim)) %>%
  mutate(cumulative_reward_lower_CI = avg_cumulative_reward - 1.96*se_cumulative_reward,
         cumulative_reward_upper_CI = avg_cumulative_reward + 1.96*se_cumulative_reward)%>%
  filter(t <= max_obs)

# TODO: Make a plot to compare the performance of the UCB policy and epsilon-greedy policy
# 1: A plot that shows only the average cumulative rewards over time using the df_history_agg dataframe
# 2: The plot as defined in (1) together with the 95% confidence interval.

```

Task 5: to better understand the  $\epsilon$ -greedy policy and UCB algorithm, plot for each the cumulative % of the time a certain arm is chosen. Which arm is chosen most often? does this differ between the two approaches?

```

# dataframe with arm choices for UCB
df_arm_choices_UCB_01 <- df_Yahoo_UCB_01%>%
  group_by(t, sim, choice, .drop=FALSE) %>%
  mutate(n = sum(!is.na(reward))) %>% # Count valid rewards
  group_by(t, choice, .drop=FALSE)%>%
  summarise(sum_n = sum(n)) %>% # Sum the counts of chosen arms
  group_by(choice, .drop=FALSE)%>%
  mutate(cum_sum_n = cumsum(sum_n), # Cumulative sum of chosen arms
         avg_cum_sum_n = cum_sum_n / (10 * t)) %>% # Average cumulative selections
  filter(t <= max_obs) # Filter by max observations

# TODO: Plot the average % of arms chosen for the UCB policy

```

```

# dataframe with arm choices for Epsilon Greedy
df_arm_choices_EPS_01 <- df_Yahoo_UCB_eps_greedy %>%
  group_by(t, sim, choice, .drop=FALSE) %>%
  mutate(n = sum(!is.na(reward))) %>% # Count valid rewards
  group_by(t, choice, .drop=FALSE)%>%
  summarise(sum_n = sum(n)) %>% # Sum the counts of chosen arms
  group_by(choice, .drop=FALSE)%>%
  mutate(cum_sum_n = cumsum(sum_n), # Cumulative sum of chosen arms
         avg_cum_sum_n = cum_sum_n / (10 * t)) %>% # Average cumulative selections
  filter(t <= max_obs) # Filter by max observations

# TODO: Plot the average % of arms chosen for the Epsilon Greedy policy

```